

# Cloud Resource Management

## Introduction

1. Cloud is multi-tenant by its very nature with QoS guarantees for different applications and different programming frameworks
  - Computational resources are CPUs, memory footprint, network bandwidth and latency and must be accounted for when scheduling resources
  - Look at state-of-the-art for managing cloud resources


## Resource Management for the Cloud

1. Context and terminologies
  - Resources in the context of cloud applications
    - CPU
    - Memory footprint and effective memory access time
    - Storage (bandwidth, latency)
    - Network (bandwidth, latency)
  - Resource utilization
    - Percent time a resource is actively used
  - An application may use all or most of the resources
  - Begs the questions
    - Which resource's utilization is more important?
    - Should we optimize for one or multiple resources' utilization?
  - Basics
    - CPU scheduling in traditional OS
      - \* Focus on CPU utilization
      - \* FCFS (first come first served), SJF (shortest job first), Round-robin, SRTF (shortest remaining time first), priority queues, multi-level priority queues...
    - Fairness
      - \* User's perception of resource allocation policy
      - \* e.g., round-robin gives a feeling of fairness if all processes are created equal; Each process gets  $1/N$  processor resource, where  $N$  = number of processes
    - What if all processes are not created equal?
      - \* "Fair sharing" takes into account priorities; you get what you pay for
  - Variants of fair sharing
    - Equal share
    - Max-min fairness
    - Weight max-min fairness
  - How do we extend "fair sharing" to the cloud?
    - Need to consider all the resources, not just CPU
  - Computations running in data centers are not all uniform
    - Data intensive workloads (health informatics)
    - Compute intensive workloads (ML algorithms)
    - Network intensive workloads (Netflix)

## Fair Share Schedulers

1. Hadoop
  - Uses max-min scheduling
    - K-slots per machine (CPU + memory)
      - \* A slot is a fraction of the machine's resources
  - A job may consist of a number of tasks
  - Assign one task per slot
  - Apply max-min fairness in mapping tasks to slots

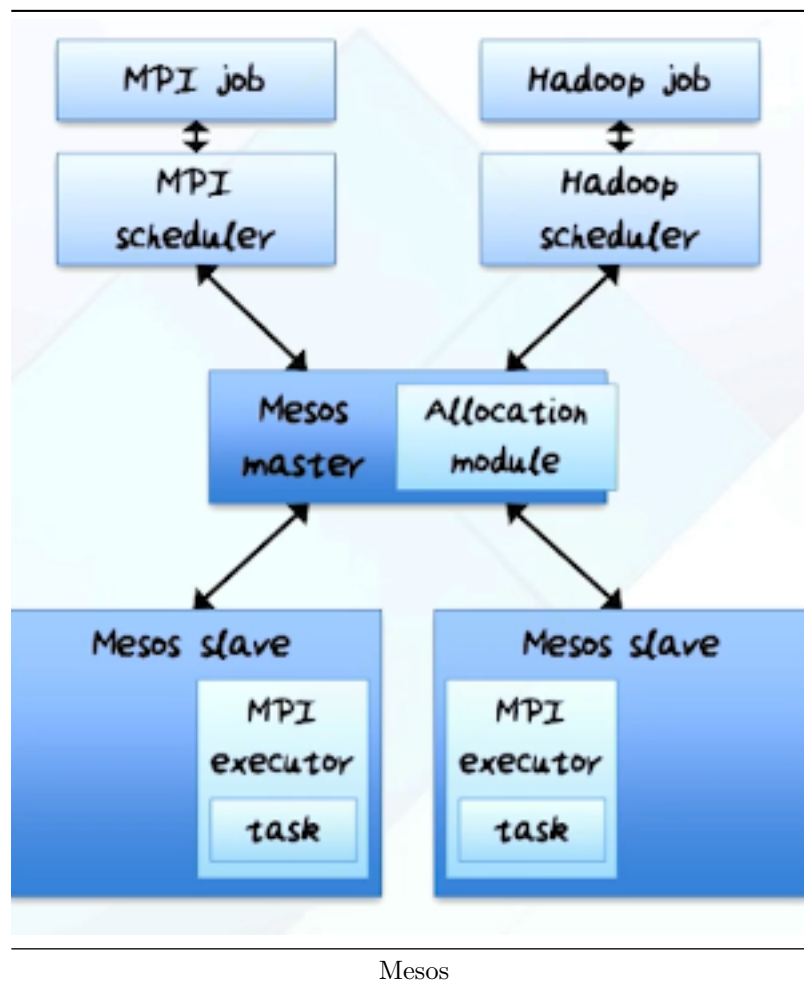
- Could result in under-utilization of slot resources or over-utilization, e.g. memory thrashing
    - If a task needs 128 MB, it will lead to thrashing in the above example
    - If it only needs 32 MB then half the allocation is wasted
2. Dominant Resource Fairness (DRF)
- Takes a holistic view of a machine's resources
    - Total resource of a machine: 10 CPUs and 4 GB
    - Task needs: 2 CPUs, 1 GB
      - \* Dominant resource of task 1 is memory
      - \* If task 1 is allocated a slot on this machine its “dominant resource share” is 25%
  - DRF uses max-min fairness to dominant shares
    - System: 9 CPUs + 18 GB
    - A: each task 1 CPU + 4 GB (dominant: memory)
    - B: each task 3 CPU + 1 GB (dominant: CPU)
    - Equalize allocation for A and B on dominant shares
      - \* Each gets 2/3 of the dominant share



## Mesos

1. Challenge for resource sharing in the cloud
  - Applications use a variety of frameworks
    - Map-reduce Dryad, Spark, Pregel
  - Apps written in different frameworks need to share the data center resources simultaneously
    - 1000s of nodes, hundreds of “jobs”, each with millions of (small duration) tasks

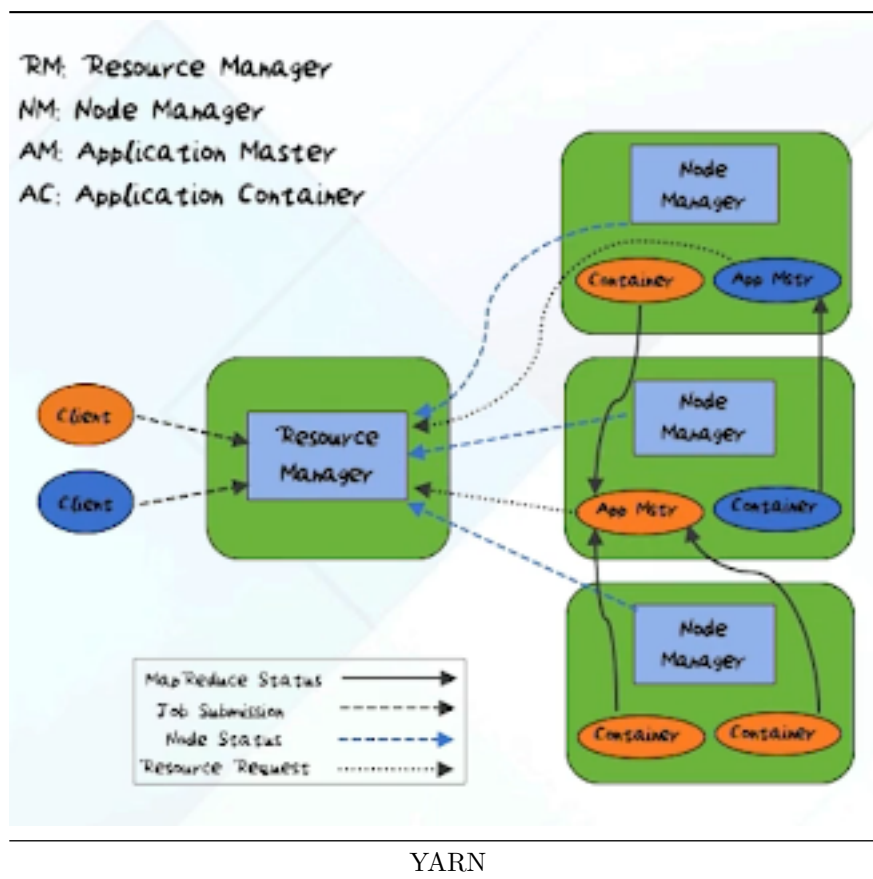
- Need to maximize utilization while meeting Qos needs of the apps
2. Traditional approach to resource sharing
    - Static partitioning (cluster with 9 machines)
      - Each framework gets 3 machines
  3. Optimal approach
    - Global schedule
      - Inputs: Requirements of frameworks, resource availability, policies (e.g., fair sharing)
      - Output: Global schedule for all the tasks
      - Not practical: Too complex
        - \* Not scalable to the size of cluster and number of tasks
        - \* Difficult to implement in the presence of failures
    - Mesos goal
      - Dynamic partitioning leading to high utilization
  4. Mesos approach: Fine-grained resource sharing
    - A thin layer (a la micro-kernel) allocation module between the scheduler of the frameworks and the physical resources



5. Mesos: Resource offers
  - Make offers to the frameworks
  - Framework assigns tasks based on the offer
  - Framework can reject the offer if it does not meet its needs
  - Solution may not be optimal but resource management decisions can be at a fine grain and quick

## Hadoop YARN

1. YARN: Yet Another Resource Navigator
  - Similar goal to Mesos
    - Resource sharing of the cluster for multiple frameworks
  - Mesos is “offer” based from allocation module to the application frameworks
  - YARN is “request” based from the frameworks to the resource manager
2. YARN background
  - Traditional Hadoop (open source implementation of map-reduce)
    - Job-tracker/Task-tracker organization
    - Poor cluster utilization
      - \* Distinct map/reduce slots
3. YARN details
  - Client submits app to resource manager (RM)
  - RM launches application master (AM) in a container
  - AM registers with RM at boot-up, requests application containers (ACs) to RM
  - RM allocates ACs in concert with the node manager (NM)
  - AM contacts the NM to launch its ACs
  - Upon application completion, AM shuts down
  - YARN can implement different scheduling policies
    - FCFS, fair, capacity



## Google Borg Resource Manager

1. Google's Resource Manager
  - Similar in spirit to Yahoo's YARN and Mesos

- Provide data center resources for a variety of applications, written in a variety of frameworks, running simultaneously
- 2. Borg terminologies
  - Cluster: A set of machines in a building
  - Site: Collections of buildings
  - Cell: A subset of machines in a single cluster
    - Unit of macro resource management
- 3. Borg priority bands
  - Production jobs
    - Higher priority jobs (e.g., long-running server jobs)
  - Non-production jobs
    - Most batch jobs (e.g., web crawler)
  - Production jobs higher priority
  - Non-overlapping “priority bands”
    - Monitoring, production, batch, best-effort
- 4. Borg architecture and scheduling
  - BorgMaster: Per cell resource manager
  - Borglet: Agent process per machine in cell
  - Scheduler
    - High-low priority jobs
    - Round robin within a band
    - Consults BorgMaster for resources
    - BorgMaster
      - \* Feasibility checking
      - \* Scoring
    - Tasks
      - \* Run in containers



- Application choice
  - Trade performance guarantees of allocation latency
- 3. Mercury architecture
  - Central scheduler
    - Policies/guarantees
    - Slower decision making
  - Distributed scheduler
    - Fast/low-latency decisions
  - AM specifies resource type
    - Guaranteed containers
    - Queueable containers



## Conclusion

1. Cloud resource management technologies are not unique from what occurs in distributed systems
  - Scale at which these algorithms must operate and QoS requirements make building such software a non-trivial engineering exercise