

# Data Center Networks

## Introduction

1. Cover the following topics:
  - Metrics for network performance evaluation
  - Tools for testing and debugging data center networks
  - Case studies of analyzing data center network traffic
2. Answer the following questions:
  - What can go wrong in data center networks?
  - What are the tools for testing and debugging?
  - What are the tools for measurements and performance analysis?
  - Case studies of measurements of data center networks?

## Data Center Networking Overview

1. Metrics
  - Latency: Time between two points
  - Throughput: Events per unit time
  - Utilization: How well the network infrastructure is used
  - Scalability: As you increase the offered load into the system, does the performance experienced by the applications remain unchanged?
    - Is my performance affected by other applications running?
2. Performance Evaluation
  - Modeling: Mathematical way of representing what's going on in the system
    - Computer systems are mostly heuristics, can be difficult
    - More approximate than a simulation
  - Simulation: Taking the real system and writing a program that simulates it. Analyze the performance of the simulator
    - Can drive simulation with traces from a real system
  - Implementation and Measurements: Build the system and study its performance
    - Expensive
  - Ideally, we first model, then simulate, then implement

## What Can Go Wrong?

1. Typical networking related problems
  - Forwarding loops
  - Link failures
  - Forwarding inconsistencies -> often leads to forwarding loops
  - Unreachable hosts
  - Unwarranted access to hosts (which should not be reachable)
2. Data Center vs Traditional Networks
  - Forwarding loops
    - In traditional networks caused by failure of spanning tree protocols
  - Link failures
    - Response is different, but problem is the same
  - Unreachable hosts
    - In traditional networks due to errors in ACLs or routing entries
    - In SDNs due to missing forwarding entries
  - Unwarranted access to hosts
    - In traditional networks due to errors in ACLs
    - In SDNs caused by unintended rule overlap
3. Challenges for Data Center Networks
  - Potential control loops

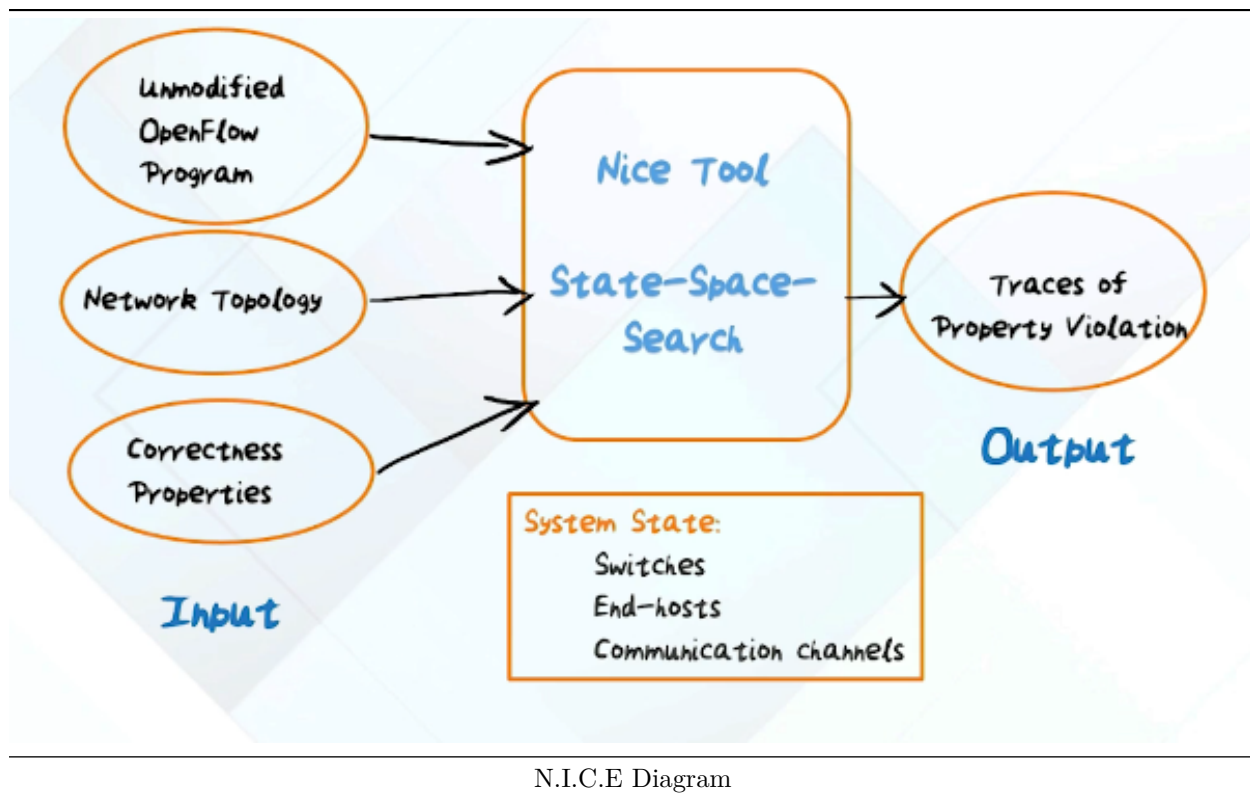
- Switches, controller, application
- End-to-end behavior of networks
  - Networks are getting larger
  - Network functionality becoming more complex
  - Partitioning functionality: Unclear boundary of functionality between network devices and external controllers
- 4. Effect of Network Bugs
  - Unauthorized entry of packets into a secured zone
  - Vulnerability of services and the infrastructure to attacks
  - Denial of critical services
  - Effect network performance and violation of SLAs

## Tools for Testing, Debugging, Verification

1. Approaches to Testing and Debugging
  - Domain-specific languages to minimize errors
    - Don't use C to program controller to limit kinds of errors
  - Limited set of principles
  - Symbolic execution and model checking
    - Constrain state space to check the model
  - Static analysis of the network state
    - Run network, observe state, do a post-mortem later on
  - Live debugging
    - As the applications are running, use a network debugger to probe

## N.I.C.E

1. Can we build a tool for systematically testing data center applications?
  - Data plane
    - Huge space of possible packets
  - Control plane
    - Huge space of event ordering on the network
  - Light at the end of the tunnel
    - Equivalence classes of data packets
    - Application knowledge of event ordering
2. N.I.C.E: No bugs In Controller Execution
  - Tool for automatically testing OpenFlow applications
    - Goal: Systematically test possible behaviors to detect bugs
  - Possible state-space is exponential
    - Combinatorial explosion with brute force approach
  - NICE's magic sauce
    - State-space exploration via Model Checking (MC)
    - Combine Symbolic Execution (SE) with Model Checking to prevent state-space explosion
  - Take advantage of knowledge of the application properties to reduce the state-space



## OFRewind

1. Static analysis of OpenFlow programs
2. How does it work?
  - In production: Record state (events and traffic) with minimal overhead
  - Later: Replay state at convenient pace
  - Troubleshoot: Reproduce problems at chosen times/locations
3. Keys to scalability of this approach
  - Record control plane traffic
  - Skip/aggregate data plane traffic
  - Replay: Best effort as opposed to deterministic
    - Identify performance problems or correctness issues, so we don't necessarily need a deterministic execution
  - Over-arching goal: Partial recording/replay of chosen times/locations to reproduce problems
4. How to use OFRewind
  - Deploy OFRecord in production
    - "Always on" OF messages, control plane, data plane summaries
    - Selection rules as necessary
  - Deploy OFRecord in the lab
    - Localize bugs and validate bug fixes

## Ndb

1. Tool for live debugging of errant network behavior (similar to gdb for program control flow)
2. Approach
  - Use SDN architecture to systematically track down network bugs
  - Capture and reconstruct the sequence of events leading to the errant behavior
  - Network breakpoints defined by the user

- Filter (header, switch) to identify the errant behavior
- Back trace generation
  - Path taken by the packet
  - State of the flow tables at each switch
- 3. How it works
  - Breakpoint is a filter on packet header
    - e.g. <switch S; IP\_src = A; IP\_dst = B; TCP\_port = 22>
    - Switches send “postcard” on matching entries to a central collector
    - Collector stores the postcards to construct a “backtrace”
  - Collector can become a bottleneck
    - Instead of applying rule to all switches, apply to a subset to reduce total amount of collected traffic
    - Collector can be parallelized, don’t have to be centralized

## Header Space Analysis and Netplumber

1. Elements of Header Space Analysis
  - General model that is agnostic to the protocols and network topologies
  - Models the packet header (length L) as a point in an L-dimensional hyperspace
  - Models all network boxes as a transformer on the packet header space
  - Defines an algebra for this transformation
    - Composable, invertible
  - Models all kinds of forwarding functionalities regardless of specific protocols and implementations
2. Using the Model
  - All traffic flows can be expressed as a series of transformations using the algebra
  - Allows asking questions such as
    - Can two hosts communicate with each other?
    - Are there forwarding loops?
    - Are there network partitions?
3. Netplumber
  - A system built on header space analysis
  - Creates a dependency graph of all forwarding rules in the network and uses it to verify policy
    - Nodes in the graph -> Forwarding rules in the network
    - Directed edges -> Next hop dependency of the forwarding rules
4. Represent forwarding policy as a dependency graph
  - Flexible policy expression
    - Probe and source nodes are flexible to place and configure
  - Incremental update
    - Only have to trace through dependency sub-graph affected by an update to the forwarding policy
  - Parallelization
    - Can partition dependency graph into clusters to minimize inter-cluster dependences

## Veriflow

1. Tackles the problem of network-wide verification of traffic flows
2. Goal: Detect routing loops, black holes, access control violations
3. Approach
  - Interpose verification layer between SDN controller and the network elements
  - Formulate network invariants from the SDN controller
  - Construct a model of network behavior
  - Monitor the network for violation of invariants

## Network Traffic Characteristics

1. Data Center Network Traffic Study
  - Are links over-subscribed?
  - Is there sufficient bisection bandwidth?
  - Is centralization (via SDN controller) feasible?
2. Setup for the Study
  - Classic model of data center network
    - Core (L3), aggregation (L2), edge (Top-of-Rack-L2) layers
  - 10 data centers from three classes
    - University, private enterprise, cloud
  - User community
    - Internal (university, private) and external (cloud)
  - Methodology
    - Analyze running applications using packet traces
    - Quantify network traffic from applications
3. Results Summary
  - Significant amount of small packets (~50% less than 200 bytes)
    - TCP acks, “I am alive” messages
  - Importance of connection persistence
  - Traffic distribution
    - Clouds: Most traffic (75%) within a rack -> good colocation of application components
    - Other DCs: 50% inter-rack -> un-optimized placement
  - Link utilization
    - Core > Aggregation > Edge
    - Bisection bandwidth sufficient (only 30% of the bisection used)
4. Insights from the Study
  - Are links over-subscribed? No
    - 75% traffic within a rack
    - Core links utilization < 25%
    - Need better load balancing, VM placement, and VM migration
  - Is there sufficient bisection bandwidth? Yes
    - Small packet sizes
    - Utilization < 30%
  - Is centralization feasible? Yes
    - Most apps use TCP/IP so setting up switches and amortizing the cost is worth it

## Classification of Traffic

1. Types of traffic
  - D2C Traffic: Traffic exchanged between Yahoo servers and clients
  - D2D Traffic: Traffic exchanged between different Yahoo servers at different locations
  - Client: Non-Yahoo host connect to Yahoo server
2. Methodology for collecting traffic data
  - Anonymized NetFlow datasets collected at the border routers of five major Yahoo data centers
    - Dallas (DAX), Washington DC (DCP), Palo Alto (PAO), Hong Kong (HK), United Kingdom (UK)
  - Meta data collected
    - Timestamp, source and destination IP address, transport layer port number, source and destination interface on the router, IP protocol, number of bytes and packets exchanged
3. Key findings of the study
  - Yahoo data centers are heirarchically structured
  - D2D traffic patterns
    - D2C triggered traffic: Smaller with higher variance commensurate with user dynamics
    - Background traffic: Dominant and not much variance

- Highly correlated traffic at data centers
  - Replicated services at different data centers
  - Implications for distributing services at multiple data centers

## Structure of Google WAN

1. Google WAN
  - Characteristics
    - Global user base
    - QoS needs: High availability and quick response
  - Implications
    - Rapid movement of large data across WAN
  - Organization of the network
    - I-scale: Internet facing
    - G-scale: Inter-data center
2. G-Scale Network
  - OpenFlow powered SDN
  - Proprietary switches from merchant silicon and open source routing stacks with OpenFlow support
  - Each site
    - Multiple switch chassis
    - Scalability (multiple terabits of bandwidth)
  - Fault tolerance
    - Sites inter-connected by G-scale network
    - Multiple OpenFlow controllers
    - No single point of failure

## Conclusion

1. Tools and techniques for testing, debugging, verification, and performance analysis of data center networks