

# Developing Virtual Network Functions

## Introduction

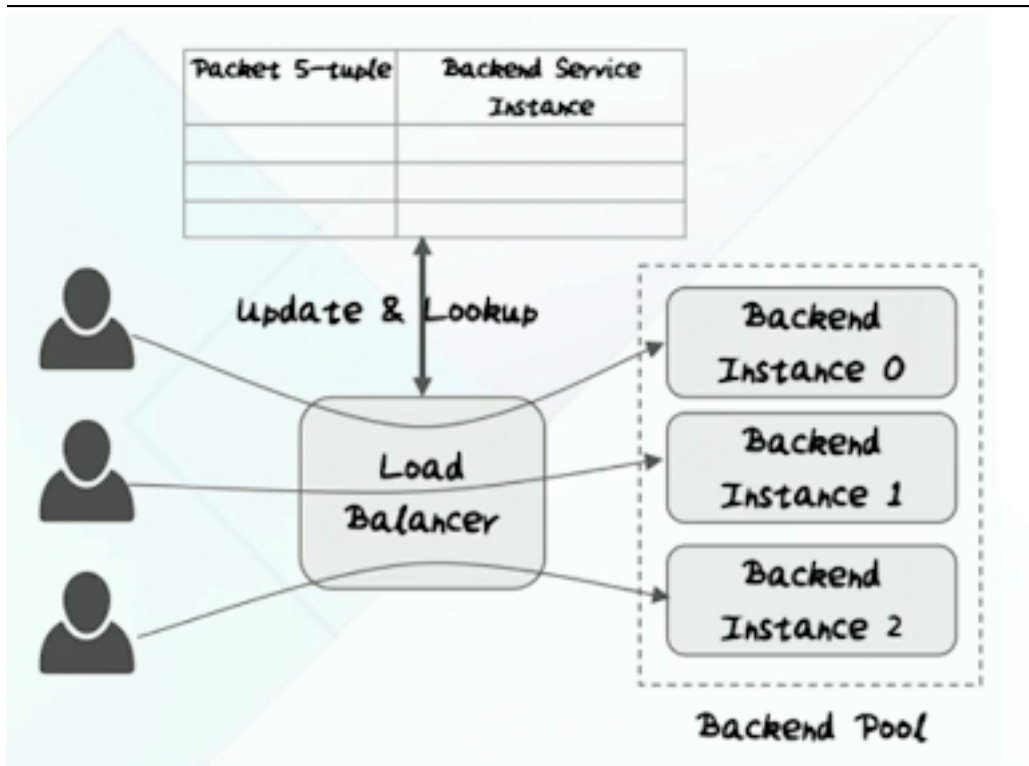
1. In-depth look at virtual network functions
  - How do we implement performance-conscious network functions?

## Developing Virtual Network Functions

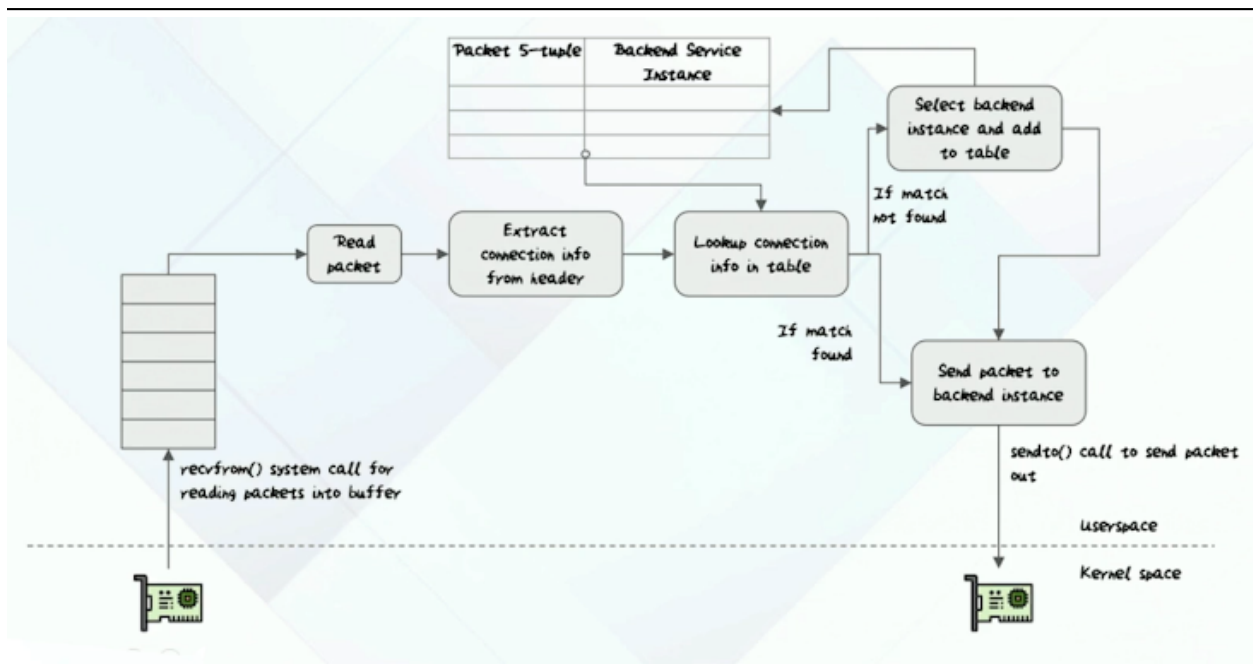
1. VNF: Virtual Network Function
  - Use concrete examples to understand how VNFs are implemented
2. DPDK: Data Plane Development Kit
  - User-space packet processing

## Virtual Network Functions

1. Overview
  - Network function implemented in user-space on top of hypervisor
    - For portability
  - Load balancer as a concrete example
    - Keeps a pool of backend service instances (e.g., HTTP server)
    - Distributes incoming packet flows to a specific instance to exploit inherent parallelism in the hardware platform and balance the load across all the service instances
2. Architecture of a load balancer network function
  - Distribute client connections to a pool of backend service instances
    - For example HTTP server
  - Use packet's 5-tuple to choose backend instance
    - Source address, destination address, source port, destination port, protocol
    - Provides connection-level affinity
    - Same connection is sent to same backend instance
  - Most network functions are implemented on top of Linux



Load Balancer Architecture



Load Balancer Diagram

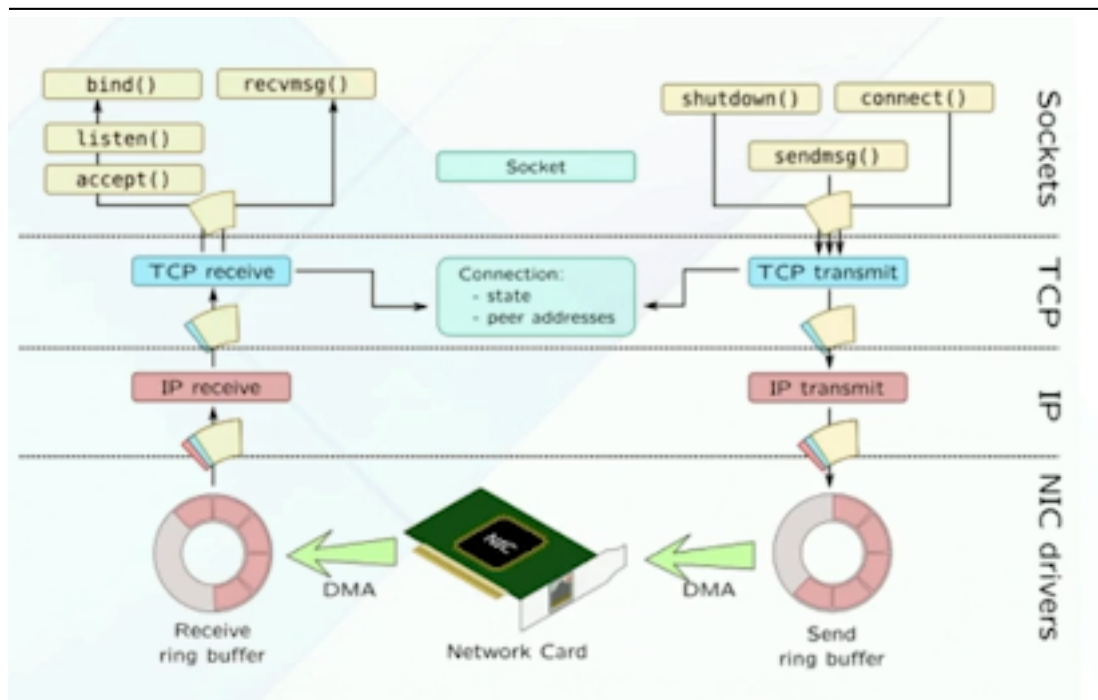
## Performance Issues in Implementing VNF

1. Eliminating the overhead of virtualization

- Network function is in the critical path of packet processing
- Need to eliminate the overhead of virtualization
  - Intel VT-d allows the NIC to bypass the VMM (hypervisor) by direct- mapping user-space buffers for DMA and passing the device interrupt directly to the VM above the VMM
- Is that enough?
  - Unfortunately no.
  - Let's look at the path of packet processing in an OS like Linux

## 2. Packet processing in Linux

- NIC uses DMA to write incoming packet to a receive ring buffer allocated to the NIC
- NIC generates an interrupt which is delivered to the OS by the CPU
- OS handles the interrupt, allocates kernel buffer and copies DMA'd packet into the kernel buffer for IP and TCP processing
- After protocol processing, packet payload is copied to application buffer (user-space) for processing by the application



Linux Packet Processing

## 3. An example networking app on Linux Kernel

- A web server on Linux
- 83% of the CPU time spent in the kernel
- This is not good news even for a networking app such as a web server
- This is REALLY BAD news if the app is a network function

## 4. Performance hits

- One interrupt for each incoming packet
- Dynamic memory allocation (packet buffer) on a per packet basis
- Interrupt service time
- Context switch to kernel and then to the application implementing the NF
- Copying packets multiple times
  - From DMA buffer to kernel buffer
  - Kernel buffer to user-space application buffer
  - Note that a NF may or may not need TCP/IP protocol stack traversal in the kernel depending on its functionality

## Performance-conscious Implementation of VNF

1. Circling back to Virtualizing Network Functions
  - Intel VT-d provides the means to bypass the hypervisor and go directly to the VM (i.e., the guest kernel which is usually Linux)
    - NF is the application on top of Linux
  - Guest kernel presents a new bottleneck
    - Specifically for NF applications
    - Sole purpose of NF applications is to read/write packets from/to NICs
  - Slowdown due to kernel is prohibitive
    - Demand scaling broke down in 2006 (leakage current started dominating)
    - CPU clock frequencies are not increasing significantly from generation to generation
    - CPU speeds are not keeping up with network speeds
    - NICs can handle more packets per second -> increasing pressure on CPU
  - So it is not sufficient to bypass the VMM for NF virtualization
    - We have to bypass the kernel as well
2. Performance-conscious Packet Processing Alternatives
  - By-passing the Linux kernel
    - Netmap, PF\_RING ZC, and Linux Foundation DPDK
  - These alternatives possess common features
    - Rely on polling to read packets instead of interrupts
    - Pre-allocate buffers for packets
    - Zero-copy packet processing: NIC uses DMA to write packets into pre-allocated application buffers
    - Process packets in batches

## Data Plane Development Kit

1. DPDK was developed by Intel in 2010
  - Now an open source project under Linux Foundation
  - Libraries to accelerate packet processing
  - Targets wide variety of CPU architectures
  - User-space packet processing to avoid overheads of Linux kernel
2. Features of DPDK
  - Buffers for storing incoming and outgoing packets in user-space memory
    - Directly accessed by the NIC DMA
  - NIC configuration registers are mapped in user-space memory
    - PICE configuration space
    - Can be modified directly by user-space application
  - Effectively bypasses the kernel for interacting with NIC
3. DPDK is a user-space library
  - Very small component in the kernel
    - Used for initialization of user-space packet processing
  - Needed to initialize the NIC to DMA to appropriate memory locations
  - Setup memory mapping for configuration registers on the NIC
    - PCI configuration space
    - Updating those registers is then done in user-space
4. Poll Mode Driver
  - Allows accessing receive (rx) and transmit (tx) queues
  - Interrupts on packet arrival are disabled
  - CPU is always busy polling for packets even if there are no packets to be received
  - Receive and transmit in batches for efficiency

```
while(true) {  
    buff <- bulk_receive(in_port)
```

```

for pkt in buff:
    out_port <- look_up(pkt_header)
    # handle failed lookup somehow
    out_buffs[out_port].append(pkt)
    for out_port in out_ports:
        bulk_transmit(out_buffs[out_port])
}

```

## NIC Ring Buffer

1. Overview
  - Each NIC queue is implemented as a ring buffer (not specific to DPDK)
  - Each slot in the ring buffer holds a “descriptor” for a packet
    - Descriptor contains a pointer to the actual packet data and other metadata
    - Actual packet is stored in another buffer data structure
  - Write pointer is advanced when NIC receives packets
  - Read pointer is advanced when CPU reads packets
2. NIC ring buffer
  - Upon packet arrival, NIC populates the next vacant slot with packet’s descriptor
  - CPU core running NF polls ring for unread slots
  - When new descriptors are found
    - CPU reads the packet data for those descriptors
    - Returns packets to application
  - No need for locking: producer and consumer are decoupled in ring buffer
  - If no vacant descriptor slots in ring buffer, NIC drops packets
3. Pre-allocated buffers for storing packets
  - Instead of allocating a buffer for each incoming packet, DPDK preallocates multiple buffers on initialization
  - Each rx queue in the NIC can hold no more packets than the capacity of the ring buffer
    - Total size of packet buffers is thereby known = capacity of ring
  - Incoming packet is DMA’d into the buffer along with adding new packet descriptor to ring buffer
  - DPDK uses huge pages to maintain large pools of memory
    - Each page is 2 MB in size (compared to traditional 4 KB pages)
    - Fewer pages -> fewer TLB misses -> improved performance
4. No overhead of copying packet data
  - NIC DMA transfers packets directly to user-space buffers
  - Protocol processing (TCP/IP) is done using those buffered packets in place
    - ...if needed by the network function (Note: not all NFs require TCP/IP processing)
5. Upshot of NF using DPDK and Intelligent NICs
  - All the kernel overheads in packet processing (alluded to earlier) mitigated/eliminated
  - Results in performance-conscious implementation of the VNF
  - Developer of NF can concentrate on just the functionality of the NF
    - DPDK alleviates all the packet processing overheads for any NF

## Implementation of VNF

1. DPDK Optimizations
  - Various optimization opportunities are available in DPDK to improve packet processing
  - Each optimization attempts to eliminate a particular source of performance drop in Linux kernel and/or exploitation of hardware features in NICs and modern CPUs
2. Implementing NFs using DPDK on commodity hardware
  - Modern commodity servers contain multi-core CPUs
  - Using multiple cores for packet processing can allow us to match the increasing capacities of NICs
  - NUMA servers

- Multiple sockets each with given nubmer of cores and local RAM
  - Accessing remote RAM is much more expensive than local RAM
- Upshot
  - Need to carefully design the packet processing path from NIC to NF taking these hardware trends into account
  - Partnership between system software and hardware
- 3. DPDK application model
  - Run-to-completion model
    - Polling for incoming packets, processing on packet, and transmission of output packet all done by the same core
    - Each packet is handled by a unique core
  - Pipelined model
    - Dedicated cores for polling and processing packets
    - Inter-core packet transfer using ring buffers
- 4. Run-to-completion model
  - All cores responsible for both I/O and packet processing
  - Simplest model
  - Each packet sees only one core
    - Works for monolithic packet processing code
    - When all the packet processing logic is contained inside a single thread
    - Simple to implement but less expressive
- 5. Pipelined execution model
  - Dedicate cores for processing NF logic
  - Some cores are dedicated for reading packets
  - Each packet sees multiple cores
    - Can be used to chain multiple packet processing logics (within an NF)
    - e.g., IN -> firewall -> router -> OUT
  - Inter-core communication done using queue buffers in memory
  - Also useful when packet processing is CPU bound, so having number of polling cores < number of processing cores is a good choice
    - Intrusion Detection System

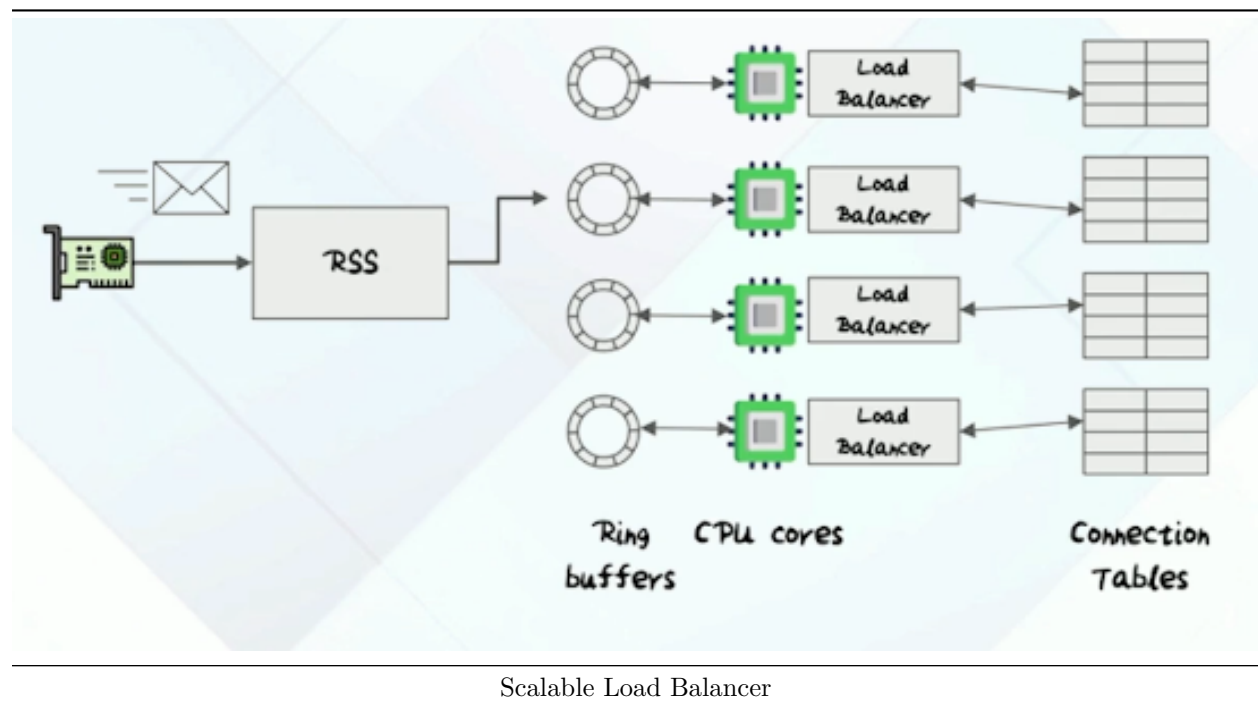
## Multi-core Implementation Challenges

1. Overview
  - How to ensure that processing done by two distinct cores don't interfere with each other?
    - If different packets of same connection are sent to different cores, sharing NF state will be a nightmare
  - How to ensure that cores participating in inter-core communication are on the same NUMA node?
  - How to ensure that the cores processing packets are on the same NUMA socket as the NIC?
2. Receive side scaling: Useful hardware technology
  - Enabler of multi-core processing
  - Use hashing to distribute incoming packets to individual cores
    - Hash function takes 5-tuple of packet as input
    - src\_ip, dst\_ip, src\_port, dst\_port, proto
  - Each core is assigned a unique ring buffer to poll
    - No contention among threads
  - Different connection -> different queue (ring) -> different core
    - Per-connection state is accessed only by a single core, so state management is easy
3. Multi-core support in DPDK
  - Allows admin to specify the following (hardware/software partnership):
  - Allows mapping of specific RX queue to specific CPU core
    - Port 0 -> Rx queue 1 -> CPU core 6
    - CPU core 6 -> Port 1 -> Tx queue 2

- Flexible to create as many queues as admin wants
  - Each thread is pinned to a specific core
    - To avoid contention
  - Each thread/core runs the same code
4. NUMA awareness in DPDK
- DPDK creates memory pools for inter-core communication on the same NUMA socket as the cores involved
  - Ring buffers are allocated on the same socket as the NIC and cores selected for processing
  - Remote memory access is minimized

## Putting it Together

1. Load balancer application
  - Multi-threaded, run-to-completion model
    - Each thread performing identical processing
  - Dedicated Rx and Tx queues for each core



2. Architecture
  - Incoming packet is directed to a particular read buffer
  - Each thread within a core is doing the following:
    - Receive a packet from the NIC
    - Read packet data for new descriptors
    - Extract connection info from header
    - Lookup connection info in table
    - If match, send packet to backend instance
    - If not match, select backend instance and add to table
  - Ring buffer logic is handled by DPDK
    - Application handles load balancing specifics

## Conclusion

1. State of the art: Implement network functions on top of Linux kernel

- General trends for accelerating packet processing to alleviate bottlenecks
  - Mechanisms for bypassing kernel
- Used DPDK as an exemplar
  - Explored how DPDK handles multi-core processors
- Enables very performance-conscious user-space implementations of network functions