

Introduction to Cloud Applications

Introduction

1. High level overview in best practices for developing scalable applications for the cloud
 - Evaluation pros and cons of migrating applications from enterprise into the cloud

Developing Large Applications for the Cloud

1. App developer's toolkit on the Cloud
 - Frameworks
 - json, j2ee, .NET, REST
 - Webserver
 - Apache MS IIS, IBM HIS
 - Appserver
 - Websphere, Weblogic
 - Database Server
 - MySQL, SQL, DB2, Oracle
 - NoSQL Databases
 - MongoDB, Hadoop
 - Development Environment
 - Eclipse, Rational, Visual Studio
 - Configuration/Deployment
 - Chef, Puppet
 - Lifecycle Management
 - CVS, Subversion, Git
 - Testing
 - LoadRunner, Rational

Best Practices for Scalable Cloud Application Development

1. Stateless design
 - Don't store state on the server so clients can communicate with any server
2. Loose coupling of app components
 - Don't share state among components
3. Asynchronous communication
 - Failures can occur, so calls should be asynchronous so the client can query the health of any calls that were made
 - Promote parallelism
4. Choice of Database
 - If application doesn't require semantics of a SQL database, a NoSQL database is likely preferable
5. Design patterns
 - Many applications fall into similar well-defined patterns, so cloud providers provide these patterns instructing how to structure a service
 - For example, a shopping cart
 - Promotes reuse of components

Best Practices for Ensuring Reliability and Availability

1. Ensuring no single point of failure
 - Redundancy in MapReduce using Zookeeper
2. Automate failure actions
 - Don't want to have to manually address every failure
3. Graceful degradation
 - Elastically increase resources, might not be able to provide service at the same level

4. Liberal use of logging
 - Understand what happened when things go wrong
 - Ensure that reliability logic is not on the critical path so it doesn't decrease performance
5. Replication
 - Key for ensuring reliability and availability
 - Partition dataset so workers can process in parallel
 - Some datasets aren't partition-able, such as Gmail inbox
 - Search is a good candidate for partitioning

Best Practices for Ensuring Security

1. Authentication
 - Are clients who they say they are?
2. Authorization
 - Identity management and access management
 - Are clients allowed to do what they're trying to do?
3. Data Security
 - Stored and in-flight data
4. Encryption and Key management
 - Important for ensuring data security
5. Auditing
 - Want to be able to understand what happened when something goes wrong

Best Practices for Ensuring Performance and Endurance

1. Performance
 - Identify metrics of importance
 - Response time, throughput, false positives, false negatives
 - Workload modeling
 - Need a representative dataset for stress testing
 - Stress testing
 - Identify bottlenecks
2. Endurance
 - Design for testing and verification
 - Built in mechanisms for independently measuring the health of components
 - Automated mechanisms for gathering and reporting health and performance statistics
 - Design for future evolution
 - Incremental addition of new features
 - Incremental testing of upgrades
 - Automated maintenance

New Normal

1. Confluence of several factors
 - Distributed computing
 - Grid computing to democratic HPC use by scientists
 - WWW
 - Gigabit networks
 - Appearance of clusters as a parallel workhorse
 - Virtualization
 - Enterprise transformation
 - Supply chain
 - B2B commerce
 - Business opportunity for computing as a utility service

2. Concrete evidence of success
 - Driving down costs and accuracy
 - FDA: Manual report to machine readable reports
 - * 99.7% accuracy
 - * Cost down from \$29 per page to \$0.25 per page
 - Data security
 - Organizations can benefit from the “sledge hammer” of Cloud service providers rather than internal IT teams
 - Spurring innovation
 - Cost of experimentation is small

Designing for Scale

1. Gap between expectation and reality
 - Expectation: Cloud resources are infinite
 - Reality: Cloud resources have finite capacity
 - Number of CPUs, network bandwidth, etc.
2. Implications
 - Design for “scale out”
 - Capacity planning is crucial
 - Carefully understand “scale unit” for your app
 - Multiple smaller deployments
 - E.g., one instance for every 100 core
 - Automate provisioning resources commensurate with the scale unit for your app
3. A real work example
 - Azure’s solution for financial risk modeling
 - Pathological example
 - Cost of insuring the entire world population
 - * Model: Stochastic analysis of the insurance cost
 - * Simulation time: 19 years on a single core
 - Azure’s solution
 - * Without changing the programming model runs on 100,000 geo-distributed cores

Challenges of Migrating to the Cloud

1. Service parameters
 - Latency, uptime/availability
 - Security
 - Legal issues
2. Enterprise services are complicated
 - N-tier (front-end, business logic, back-end)
 - Replication and load balancing
 - Intra-org firewalls
3. Solution
 - Hybrid implementation
 - Private plus public
 - Construct a model of cost/benefit of migrating each component

Conclusion

1. Cloud is not a panacea for instant scalability and performance of enterprise applications
 - Migrating and architecting an app for the cloud requires considerable thought
 - Always keep scale up as the fundamental design principle
 - Proliferation of cloud resources is a testament to cloud computing becoming the new normal