

Introduction to Network Function Virtualization

Introduction

1. Decouple network services needed for any enterprise, such as a firewall or malware inspection, from proprietary hardware appliances often referred to as “middleboxes”
 - Services can be run as software entities on standardized services using virtualization
 - Network services become virtual network functions

NFV Overview

1. Introduction to types of network functions and the pathway to virtualizing them
2. A concrete example of a technology enabler for virtualizing network functions
3. Synergy between SDN and NFV
4. NFV on public cloud

What are Network Functions?

1. Types of network functions
 - Firewall
 - Filters traffic based on pre-defined rules
 - Rules are simple since filtering is in the critical path of packet flow
 - Intrusion detection/prevention
 - Perform more complicated analysis of packet traffic
 - Identify complex patterns for network traffic belonging to an attack or suspicious activity
 - Network Address Translation (NAT)
 - Translates private IP address space to public IP address space and vice versa
 - Useful for organizations that have limited public IP network presence
 - WAN optimizers
 - Reduce WAN bandwidth consumption of an enterprise
 - Perform multiple techniques like caching, traffic compression, etc for reducing traffic and latency
 - Load balancer
 - Distribute traffic to a pool of backend services
 - Virtual Private Network (VPN) Gateway
 - Provides abstraction of same IP address space for network that are physically separate
 - Multiple sites communicate over WAN using tunnels between gateways
2. Why do enterprises need network functions?
 - Users view enterprises (Google, Amazon) as a monolith
 - Internal view of an enterprise computing environment
 - Clusters of machines serving many internal functions (sales, marketing, inventory, purchasing)
 - Employees access these clusters on-premise and remotely
 - Enterprises may have several points of presence
 - Interconnected by WAN
 - Enterprises may inter-operate
 - Internet connects Intel, Samsung, Microsoft, etc.
3. Network functions give the necessary safeguards and facilities for enterprises
 - Intrusion prevention: Performs inspection of packet payload to identify suspicious traffic
 - Firewall: Filters packets based on their src, dst IPs, ports and protocol
 - Load balancer: Evenly distributes incoming connections to one of the backend servers
 - WAN Accelerator: Reduces WAN bandwidth consumption by data deduplication and compression
 - VPN: Provides illusion of same network address space across multiple sites and provides encryption for inter-site traffic

Middleboxes

1. Middleboxes: Standalone hardware boxes (aka network appliances) providing specific network functions (e.g., firewall)
2. Consider the example of a retail organization (like Walmart) that holds inventory information on premises, but uses an enterprise datacenter for long-running batch processing (demand prediction, etc.)
 - End-clients communicate with on-premise application
 - Needs to scale horizontally to handle peak traffic -> need for load balancer
 - Limit ports for traffic -> need for firewall
 - Detect/prevent suspicious activity -> need for intrusion prevention
 - Communication with enterprise datacenter
 - Need for VPN for encryption of traffic and illusion of continuous IP address space
 - Need for WAN accelerator to reduce WAN bandwidth usage (reduce \$)
 - Office personnel access content on the Internet
 - Firewall checks and filters the websites accessed (and blocks restricted ones)
 - WAN accelerator contains HTTP proxy that can cache content (reducing WAN bandwidth and \$)
3. Middleboxes (network appliances)
 - Computer networking devices that analyze/modify packets
 - For purposes other than packet forwarding
 - Typically implemented as specialized hardware components

Examples of Middleboxes

1. Intrusion prevention system (IPS)
 - Security appliance
 - Monitors all open connections to detect and block suspicious traffic
 - Sysadmin configures signatures in IPS box to detect suspicious traffic
 - Can work in inline mode (can filter out suspicious traffic) or passive mode (analyzes packets outside critical/data path)
 - Table shows the various traffic signatures that Cisco's IPSs are pre-configured with
 - The system administrator can select all or any of these signatures to be searched for in packet traffic
 - This particular screenshot is for a search result for "botnet" showing 10 signatures that characterize botnet traffic



Signature ID	Signature Name	Latest Release Date	Alarm Severity	Release
21806/0	Swizzor Botnet Traffic	2017 Feb 03	High	5965
22142/0	SQL Botnet User-Agent uCp	2017 Feb 03	High	5965
11025/0	IRC OCC File Transfer	2017 Feb 03	Informational	5965
6537/0	Kraken Botnet Traffic	2017 Feb 03	High	5965
6537/1	Kraken Botnet Traffic	2017 Feb 03	High	5965
16754/0	PushDo Botnet	2017 Feb 03	High	5965
5448/0	IRC Bot Activity	2017 Feb 03	Low	5965
9013/0	IRCBOT_JK DNS Lookup	2017 Feb 03	High	5965
9013/1	IRCBOT_JK DNS Lookup	2017 Feb 03	High	5965

Botnet Traffic

2. HTTP Proxy

- Performance-improving appliance
 - Caches web content to reduce page-load time
 - Reduces bandwidth consumption
 - Can filter out blocked websites
3. Middleboxes in core cellular networks
 - Serving gateway (S-GW)
 - Responsible for routing/forwarding of packets
 - Executes handoff between neighboring base stations
 - Packet gateway (P-GW)
 - Acts as interface between cellular network and Internet
 - NAT between internal IP subnet and Internet
 - Traffic shaping
 - Mobility Management Entity (MME)
 - Key control node of LTE (Long Term Evolution)
 - Performs selection of S-GW and P-GW
 - Sets up connection when device is roaming
 - Home Subscriber Server (HSS)
 - User identification and addressing using IMSI (International Mobility Subscriber Identifier) number
 - User profile info: service subscription rates and QoS
 4. How are middleboxes different from routers and switches?
 - Middleboxes are stateful
 - Packet processing is dependent on fine-grained state
 - Updated frequently (per packet/per connection)
 - Middleboxes perform complex and varied operations on packets

Network Management and Proliferation of Middleboxes

1. Similar challenges that motivated shift of IT to cloud services
2. Leads to lock-in to the hardware vendor of each specific middlebox
 - Difficult and expensive to migrate to a different solution
3. Failures of middleboxes lead to network outages
4. High capital and operational expenditure
 - Provisioning is done based on peak capacity
 - Management/maintenance cost is high

Network Services as Software Entities

1. Network functions as software entities on COTS servers
 - Replace middleboxes with software entities
 - Run such network functions as an “application” on general-purpose servers
 - Benefits
 - Low cost of deployment
 - Better resource utilization
 - Scaling is easily possible: lower capital expense
 - Can switch between vendors easily
 - Failures are easier to deal with
2. Examples of “software” middleboxes
 - Linux iptables: provides NAT and firewall
 - SoftEther VPN
 - Squid HTTP proxy
 - nginx load balancers
 - Bro Intrusion Detection System (1999)
3. Fundamental components of software middleboxes

- Use Unix sockets -> opening a socket creates a file descriptor
 - Use system calls read() and write() to Linux kernel for reading and writing to a socket
 - Raw Linux sockets enable developer to read/write raw bytes (MAC layer data) from/to NIC
4. Architecture of a load balancer network function
- Distribute client connections to a pool of backend service instances
 - For example, HTTP server
 - Use packet's 5-tuple to choose backend instance
 - 5-tuple: src address, dst address, src port, dst port, protocol
 - Provides connection-level affinity
 - Same connection is sent to same backend instance
 - What happens when a packet arrives?
 - NIC uses DMA to write incoming packet to memory
 - NIC generates an interrupt
 - CPU handles the interrupt, allocates kernel buffer and copies DMA'd packet into buffer for IP and TCP processing
 - After protocol processing, packet payload is copied to application buffer (user-space)

Virtualization Technology

1. Why virtualization for NF?
 - Using a VM for hosting NF (instead of running NF on bare metal servers)
 - Better portability because entire environment can be deployed, all dependencies are inside VM image
 - Network management becomes easier
 - Each NF instance is shielded from software faults from other network services
2. How to virtualize?
 - Traditionally two approaches
 - Full virtualization
 - Para virtualization
 - Full virtualization is attractive since the VM on top of hypervisor can run unmodified
 - “Trap-and-emulate” technique in the hypervisor to carry out privileged operations of the VM which is running in user mode
 - Unfortunately, for network functions that are in the critical path of packet processing this is bad news
3. How “Trap-and-emulate” works
 - I/O is performed via system calls (privileged operation)
 - When guest VM performs I/O operation
 - Executes system call
 - Guest kernel is context switched in
 - Privileged instructions are invoked for reading/writing to I/O device
 - But guest kernel is actually running in user-space!
 - Guest VM is a user-space program from the host's perspective
 - Execution of privileged instruction by user-space program results in a trap
 - Trap is caught by the hypervisor
 - Performs the I/O on behalf of the guest VM
 - Notifies the guest VM after I/O operation finishes
4. Downsides of “Trap-and-emulate” for NF
 - Host kernel (e.g., Dom-0 in Xen) has to be context switched in by the hypervisor to activate the network device driver and access the hardware NIC
 - Duplication of work by the virtual device driver in the guest and the actual device driver in the host
 - NF incurs the above overheads
 - For each packet that is sent to the NIC
 - For each packet received from the NIC

- NF is in the critical path of network processing and such overheads are untenable

Eliminating Overhead of NFV

1. How do we solve the overhead problem?
 - Fortunately, hardware vendors have been paying attention
 - We will mention two approaches to eliminate I/O virtualization overheads
 - Intel VT-d
 - Intel SR-IOV
2. Intel Virtualization Technology for Directed I/O (VT-d)
 - Allows efficient access to host I/O devices (e.g., NIC)
 - Avoids overheads of trap-emulate for every I/O access
 - Allows remapping of DMA regions to guest physical memory
 - Allows interrupt remapping to guest's interrupt handles
 - Effectively direct access for guest machine to I/O device hardware
 - Configuration registers are mapped to guest VM's memory for memory-mapped I/O
3. Benefits of VT-d
 - Avoid overheads of trap-and-emulate
 - DMA by NIC is performed to/from memory belonging to guest VM's buffers
 - Interrupts are handled directly by the guest instead of hypervisor
 - Effectively, the NIC is owned by the guest VM
4. Single Root I/O Virtualization (SR-IOV) Interface
 - An extension to the PCIe specification
 - Each PCIe device (physical function) is presented as a collection of virtual functions
 - Practical deployments have 64 VFs per PF
 - Each virtual function can be assigned to a VM
 - Allows higher multi-tenancy and performance isolation
 - Separate config register space for each VF
5. Benefits of SR-IOV
 - Allows same physical NIC to be shared by multiple guest VMs without conflicts

Putting it All Together

1. Virtual Network Function implementation
 - Host machine (NICs, etc.)
 - SR-IOV
 - VT-d direct access
 - Virtual machine with DPDK driver
 - DPDK covered in the next lecture
 - NFs implemented as a user-space application running inside VM
 - DMA from SR-IOV VF directly into VM buffers

Conclusion

1. Hardware vendors made middleboxes to handle specialized network functions
 - Middleboxes sat between enterprise computing and wide-area Internet
 - This became a network-management nightmare and would make cloud computing an impossibility
 - To ensure these functions run in a platform-agnostic manner, it makes sense to execute them in a virtual layer
 - Intel developed special hardware to eliminate the overhead inherent to virtualizing these functions