

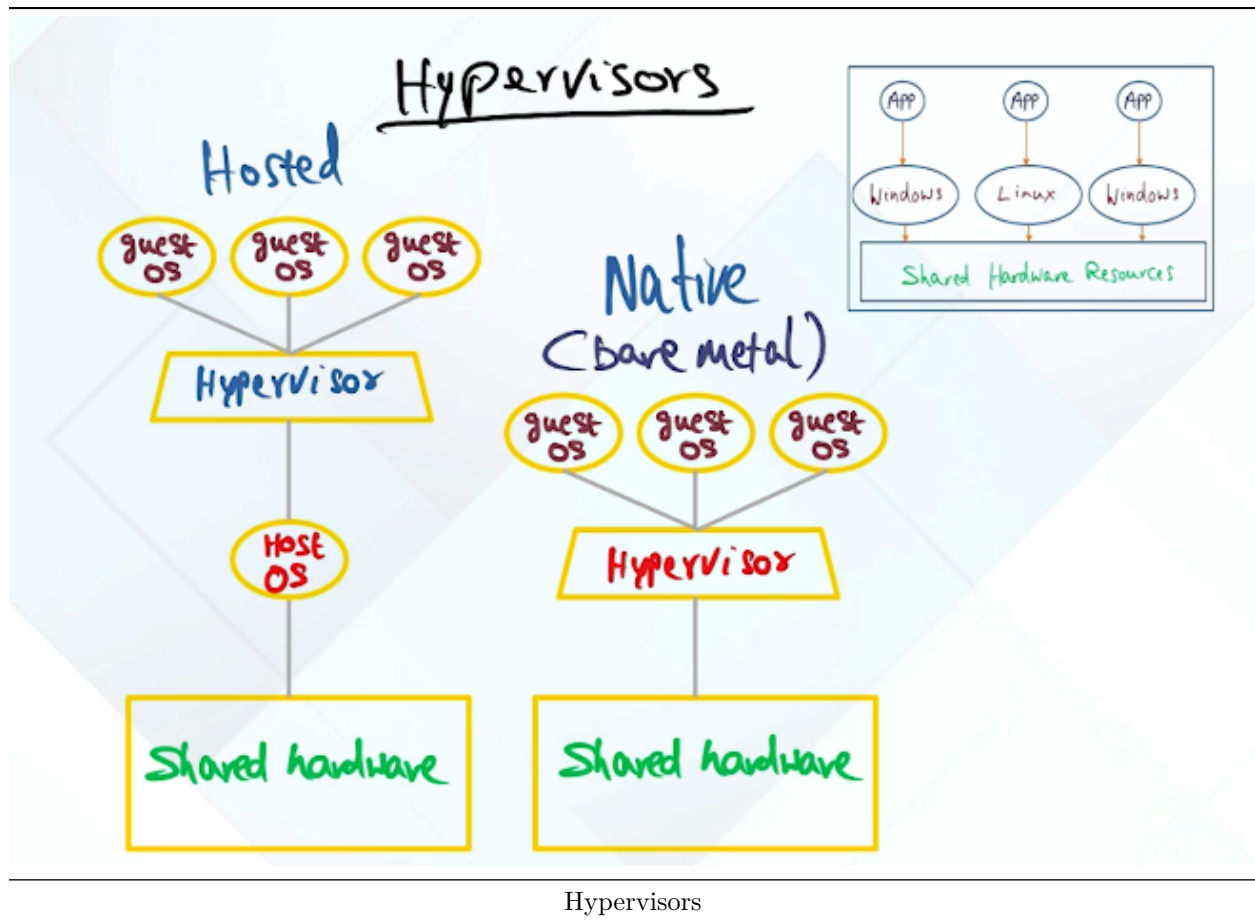
Introduction to Virtualization Technologies

Introduction

1. Virtualization dates back to IBM systems in the 60s and 70s
 - Resurrected in the 80s and 90s at universities due to a need for resource consolidation
 - Taken a new form under cloud computing


Virtualization Technologies

1. Origins of Virtualization
 - IBM VM 370 (70s)
 - Microkernels (late 80 and 90s)
 - Extensibility of OS (90s)
 - SIMOS (late 90s)
 - Simulate processor design in the context of the OS
 - Xen and VMWare (early 2000s)
 - Now...
 - Accountability for usage in data centers
2. Hypervisors
 - Run multiple OS on same hardware resources
 - Hypervisor: Thin layer that sits between guest OS and hardware resources to moderate the usage of these resources
 - Bare metal hypervisor: Running directly on bare metal
 - Hosted hypervisor: Hypervisor runs as a process on top of the host OS
 - Full virtualization: Hypervisor runs an unchanged OS binary
 - Pro: Vendors don't have to change their software
 - Con: Performance disadvantages
 - Para virtualization: Hypervisor runs a modified OS binary that knows it is virtualized (modifications 1-2% of codebase)
 - Pro: Can alleviate performance concerns
 - Con: Requires vendors to participate
3. Big Picture
 - What needs to be done?
 - Virtualize hardware
 - * Memory
 - * CPU
 - * Devices
 - Effect data and control between guests and hypervisor




Memory Virtualization

1. Memory manager
 - Every OS supports virtual memory for processes running in the OS
 - Physical memory backs up the virtual memory
 - In a virtualized context, physical memory is also an illusion
 - “Machine memory” in a virtualized world
 - Machine memory is contiguous, physical memory is not
 - OS still wants the illusion of contiguous memory
2. Dynamically increasing memory
 - Guest OS might need more physical memory to handle processes that are currently running within it
 - Guest OS goes to the hypervisor and requests more memory
 - If hypervisor has machine memory, it can give it to the guest OS, but it is also possible that there is no additional memory available
 - * Can't necessarily take memory from another OS



Ballooning

1. Understanding between hypervisor and guest OS in either full or para virtualized environment
 - Use a technique called ballooning to manage memory
 - Guest OS installs a device driver to run on the hypervisor
 - Device driver inflates the balloon when the hypervisor needs more memory
 - This returns memory to the hypervisor, which can give it to the other guest OS
 - * Some data may need to be paged out to disk
 - Can also deflate the balloon when the hypervisor needs less memory
 - * Can page in from disk
 - Balloon driver assumes cooperation between hypervisor and guest OS
 - Hypervisor has policies to determine where to take memory from




Sharing Memory Across VMs

1. The same running process in two different VMs can share memory
 - Page tables are internal to guest OS, of which the hypervisor has no knowledge
2. VM Oblivious Page Sharing
 - Look at page coming from disk and hash the contents
 - Hypervisor maintains a hint frame (hash table)
 - Hash
 - VM
 - PPN
 - MPN
 - If a page fault occurs and the hash matches a hint frame, we might be able to map it to physical memory that already exists
 - Do a full comparison of the pages
 - Mark page entries in VM as copy on write to prevent sharing modified pages
 - Don't have to do anything with respect to the VMs themselves
 - Scanning pages takes some cycles, so it's done as a background activity on the server

Successful match

- modify PPN → MPN for VM2
- mark as Copy on Write



Free up machine page

Scanning Pages

- background activity of server

Sharing Memory Across VMs

Memory Allocation Policies

1. Pure share based approach
2. Working-set based approach
3. Dynamic idle-adjusted shares approach
 - Tax idle pages more than active pages
4. Reclaim most idle memory
 - Allow for sudden working set increases

CPU Virtualization

1. Hypervisor wants to provide the illusion that each guest OS owns the CPU despite it being a shared hardware resource
2. Illusion of ownership of CPU for each VM
 - Proportional share: Proportional to the negotiated SLAs
 - Fair share
3. Delivering events to parent guest OS
 - Address translation happens on every memory access
 - Process on guest OS makes a system call
 - Page fault occurs
 - Exception (divide by zero, segmentation fault)
 - External interrupts
 - Events are delivered as software interrupts
 - In a para virtualized environment, hypervisor can provide APIs that the guest OS can use

- * Guest OS knows it isn't in control of the hardware
- If guest OS wants to switch from user mode to kernel mode, it can't because, as far as the hypervisor is concerned, it is a process running in user mode
 - * Trap into the hypervisor, hypervisor can emulate what the guest OS wants to do

Device Virtualization

1. Full virtualization
 - Trap and emulate: Guest OS thinks it has privilege to do whatever it wants, but it doesn't. Traps into hypervisor which emulates the event
2. Para virtualization
 - More opportunity for virtualization
 - Can make specific requests to the hypervisor
 - Shared buffers between guest OS and hypervisor

Control Transfer

1. Full virtualization
 - Implicit (traps) guest OS -> hypervisor
 - Software interrupts (events) hypervisor -> guest
2. Para virtualization
 - Explicit (hypercalls) guest OS -> hypervisor
 - Software interrupts (events) hypervisor -> guest
 - Guest has control via hypercalls on when event notifications are delivered
 - * Can improve performance

Data Transfer

1. Full virtualization
 - Data transfer is implicit
 - In OS, an I/O transfer might need to be copied between user and kernel space
 - Guest OS thinks it can do this itself, but it can't, so it traps
2. Para virtualization
 - Data transfer is explicit -> opportunity to innovate
 - Xen's asynchronous I/O rings are used to move data between the guest OS and hypervisor
 - Request producer and consumer hold separate pointers to the same I/O ring buffer
 - Response producer and consumer are also separate pointers

Containers

1. Containers
 - Light form of resource virtualization
 - In contrast to a full blown virtual machine of a hypervisor
 - Multiple containers on top of the same kernel
 - Illusion to each container that they are the only one using the hardware resources
 - Kernel services (e.g., I/O subsystems) not replicated in each container
2. Advantages of Containers
 - Fault isolation
 - No leaks across containers
 - Resource isolation
 - Performance guarantees for each container
 - Security isolation
 - Configuration and namespace independence
3. Containers vs VMs
 - Performance

- Order of magnitude faster to boot up container compared to spinning up VMs
- More manageable
 - Fewer OSes to manage
 - Easier patching of OS and apps
 - Improved resource utilization

VM Migration

1. VM Migration
 - Why?
 - Rebalancing resources
 - Upgrades
 - Challenges
 - Minimizing downtime
 - Avoiding disruption of active services
 - Migration options
 - Push phase
 - Stop and copy phase
 - Pull phase
2. Generic steps in migration
 - Pre-migration
 - Reservation
 - Iterative pre-copy
 - Reduce the downtime
 - Stop and copy
 - Suspend VM on source machine
 - Set up routing redirection
 - Sync all remaining VM state
 - Commitment
 - Release source machine resources
 - Activation
 - Start VM on destination resources

Conclusion

1. Technology pertaining to virtualizing data center resources is more involved than can be described in a single lecture