

# Trending Cloud Infrastructure

## Introduction

1. Large-scale deployments of IoT platforms is causing the potential disruption of the cloud ecosystem
  - This lecture covers geo-distributed computing infrastructure

## Fog Computing

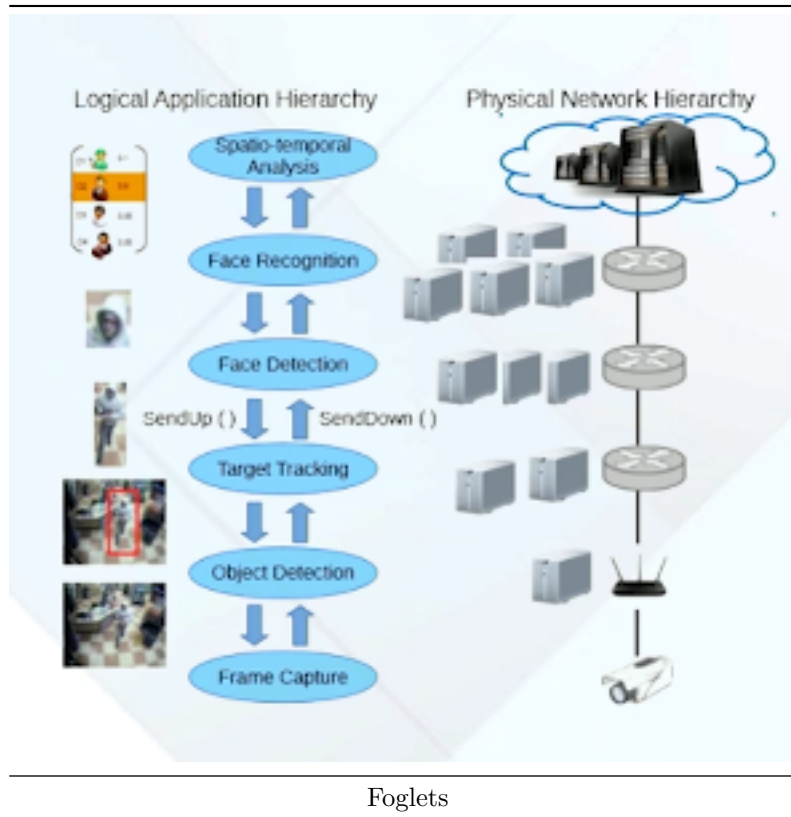
1. Realities of IoT
  - IoT devices are heterogeneous (sensing modality, data rate, fidelity, ...)
  - IoT devices are not ultra reliable
  - IoT testbed multi-tenancy
  - IoT apps are latency sensitive (sending -> processing -> actuation)
  - Need of the hour...
    - System software support for situation awareness in large-scale distributed IoT testbed
  - Requirements for system software infrastructure
    - Support multiple sensing modalities
    - Dynamic resource provisioning
    - QoS for the apps in the presence of sensor mobility
2. Limitations of Existing Cloud (PaaS)
  - Based on large data centers
    - High latency/poor bandwidth for data-intensive apps
  - API designed for traditional web applications
    - Not suitable for the future of Internet apps
3. Why?
  - IoT and many new apps need interactive response at computational perception speeds
    - Sense -> process -> actuate
  - Sensors are geo-distributed
    - Latency to the cloud is a limitation
    - Besides, uninteresting sensor streams should be quenched at the source
4. Future Internet Applications on IoT
  - Common characteristics
    - Dealing with real-world data streams
    - Real-time interaction among mobile devices
    - Wide-area analytics
  - Requirements
    - Dynamic scalability
    - Low-latency communication
    - Efficient in-network processing
5. Fog Computing
  - New computing paradigm proposed by Cisco
    - Extending the cloud utility computing to the edge
    - Provide utility computing using resources that are
      - \* Hierarchical
      - \* Geo-distributed
6. Pros of Fog Computing
  - Low latency and location awareness
  - Filter at the edge and send only relevant data to the cloud
  - Geo-distribution of compute, network, and storage to aid sensor processing
  - Alleviates bandwidth bottleneck
7. Challenges with Fog Computing
  - Manage resources
  - Deployment
  - Program the system

- Bursty resource requirements

## Programming System Exemplar

### 1. What is Foglets?

- Automatically discovers nodes
- Elastically deploy resources
- Collocate applications
- Communication and storage API
- Resource adaptation



### 2. Foglets

- PaaS programming model on the Internet of Things
- Design goals
  - Simplicity: Minimal interface with a single code base
  - Scalability: Allows dynamic scaling
  - Context-awareness: Network-, location-, resource-, and capability-awareness
- Assumes Fog computing infrastructure
  - Infrastructure nodes are placed in the fog
  - IaaS interface for utility computing

### 3. Foglets - Application Model

- Foglets application consists of distributed processes connected in a hierarchy
- Each process covers a specific geographical region

### 4. Foglets API

- Start\_App()
- On\_child\_leave()
- On\_new\_parent()
- On\_new\_child()

5. Discovery and Deploy Protocol
  - Registry/Discovery server finds all foglets near it
6. Foglets API (Communication)
  - void OnSendUp(message msg)
  - void OnSendDown(message msg)
  - void OnReceiveFrom(message msg)
  - void OnMigrationStart(message msg)
  - void OnMigrationFinish(message msg)
7. Foglets API (Context-awareness)
  - query\_location()
  - query\_level()
  - query\_capacity()
  - query\_resource()
8. Foglets - Spatio-temporal Object Store
  - App context object is tagged by key, location, and time
  - get\_object(key, location, time)
  - put\_object(key, location, time)
  - Context objects are migrated when scaling
9. Foglets - Scalability
  - Application scales at runtime based on the workload
  - Developer specifies scaling policy for each level
  - Load balancing based on geo-locations
10. Foglets Framework
  - Alleviates pain points for domain experts in mobile sensing applications (e.g., vehicle tracking, self-driving cars, etc.)
    - QoS sensitive placement of application components at different levels of the computational continuum from the edge to the cloud
    - Multi-tenancy on the edge nodes using Docker containers
    - Dynamic resource provisioning commensurate with sensor mobility
      - \* Discover and deploy, join, migration protocols

## Jetstream

1. Jetstream System
  - Streaming support
    - Aggregation and degradation as first class primitives
  - Storage processing at the edge
    - Maximize “goodput” via aggregation and degradation primitives
    - Goodput: Throughput that meets QoS metrics
  - Allows tuning quality of analytics commensurate with available backhaul bandwidth
    - Aggregate data at the edge when you can
    - Degrade data quality if aggregation not possible
2. Jetstream Storage Abstraction
  - Updatable
    - Stored data += new data
  - Reducible with predictable loss of fidelity
    - DATA -> data
  - Mergeable
    - Data + Data = Merged Data
  - Degradable
    - Local data -> dataflow operators -> approximate data

## Iridium

1. Iridium Features (Microsoft)
  - Analytics framework spans distributed sites and the core
  - Assumptions
    - Core is congestion free
    - Bottlenecks between the edge sites and the core
    - Heterogeneity of uplink/downlink edge from/to core
2. Problem being solved by Iridium
  - Given a dataflow graph of tasks and data sources
    - Where to place the tasks?
      - \* Destination of the network transfers
    - Where to place the data sources?
      - \* Sources of network transfers
  - Approach
    - \* Jointly optimize data and task placement via greedy heuristic
    - \* Goal: Minimize longest transfer of all links
    - \* Constraints: Link bandwidths, query arrival rate, etc.

## Conclusion

1. IoT is a likely disruptor of the cloud infrastructure
  - Requires a rethink of the software infrastructure of the geo-distributed computation continuum from the edge to the core
  - Programming models, storage structure, and analytics